

A thick black L-shaped frame is positioned on the left and right sides of the slide, framing the central text.

CONVOLUTIONAL NEURAL NETWORKS

By: Mohamed Aziz Tousli

Convolution

- Computer vision problems: Image classification, Object detection, Neural style transfer, Edge detection...

Vertical edge detection: $\begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}, \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}, \begin{bmatrix} 3 & 0 & -3 \\ 10 & 0 & -10 \\ 3 & 0 & -3 \end{bmatrix}$

Horizontal edge detection: $\begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

- Detect edges in general: $\begin{bmatrix} w_1 & w_2 & w_3 \\ w_4 & w_5 & w_6 \\ w_7 & w_8 & w_9 \end{bmatrix} \rightarrow$ Parameters to learn

\rightarrow Shrinking output + Throwing away information from edge \rightarrow Solution: **Padding**

- **Valid convolution**: No padding ($p=0$)


- **Same convolution**: Padding so that output size is same as input size $p = \frac{f-1}{2}$

- Strided convolution**: Convolution with pace


- $(n, n) * (f, f) \rightarrow \left(\left\lceil \frac{n+2p-f}{s} + 1 \right\rceil, \left\lceil \frac{n+2p-f}{s} + 1 \right\rceil \right)$

PS: Cross-correlation : Convolution with rotating the filter



PS: Convolution can be done in 1D and 3D the same way as in 2D



0	0	0	0	0	0	0
0						0
0						0
0						0
0						0
0						0
0	0	0	0	0	0	0



Convolution Layer



$(n, n, n_c) * (f, f, n_c) \rightarrow (n - f + 1, n - f + 1, n'_c)$; $n_c = \text{number of channels}$; $n'_c = \text{number of filters}$

- $a^{[1]} = g(w^{[1]} \cdot a^{[0]} + b^{[1]})$; $w = \text{filters}$, $a^{[0]} = \text{input}$, $b = \text{bias}$
- Filter size: $f^{[l]}$, Padding: $p^{[l]}$, Stride: $s^{[l]}$, Number of filters: $n_c^{[l]}$
- Filter: $(f^{[l]}, f^{[l]}, n_c^{[l-1]})$, Activation: $(n_H^{[l]}, n_W^{[l]}, n_c^{[l]})$, Weights: $(f^{[l]}, f^{[l]}, n_c^{[l-1]}, n_c^{[l]})$, Bias: $(1, 1, 1, n_c^{[l]})$
- Input: $(n_H^{[l-1]}, n_W^{[l-1]}, n_c^{[l-1]})$, Output: $(n_H^{[l]}, n_W^{[l]}, n_c^{[l]})$, $n_{H/W}^{[l]} = \left\lceil \frac{n_{H/W}^{[l-1]} + 2p^{[l]} - f^{[l]}}{s^{[l]}} + 1 \right\rceil$

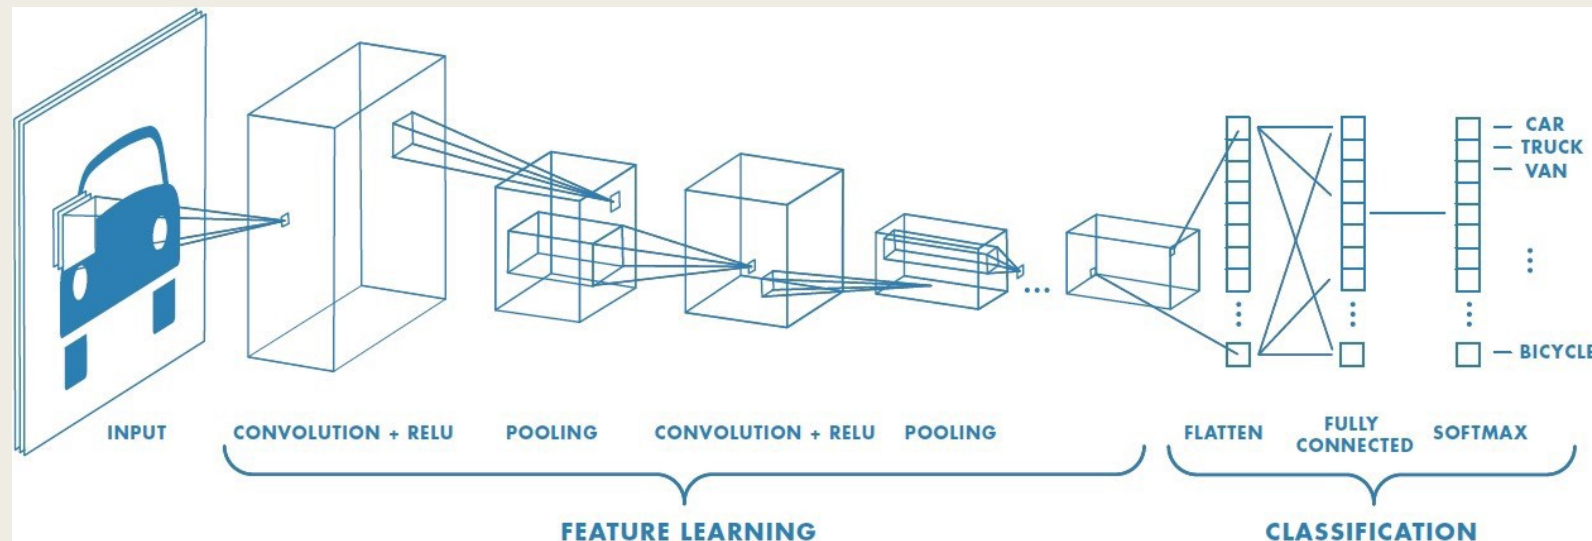
Convolutional Network

- Types of a layer in a CN: Convolution (Conv), Pooling (Pool), Fully connected (FC)
- **Max pooling:** (Hyper parameters: f (filter size), s (stride))
- Input size: $(n_H, n_W, n_c) \rightarrow \left(\left\lceil \frac{n_H - f}{s} + 1 \right\rceil, \left\lceil \frac{n_W - f}{s} + 1 \right\rceil, n_c \right)$

2	1	3	1
1	0	1	4
0	6	9	5
7	1	4	1



2	4
7	9

PS: There is also **average** pooling + There is no padding in pooling + There are no parameters to learn

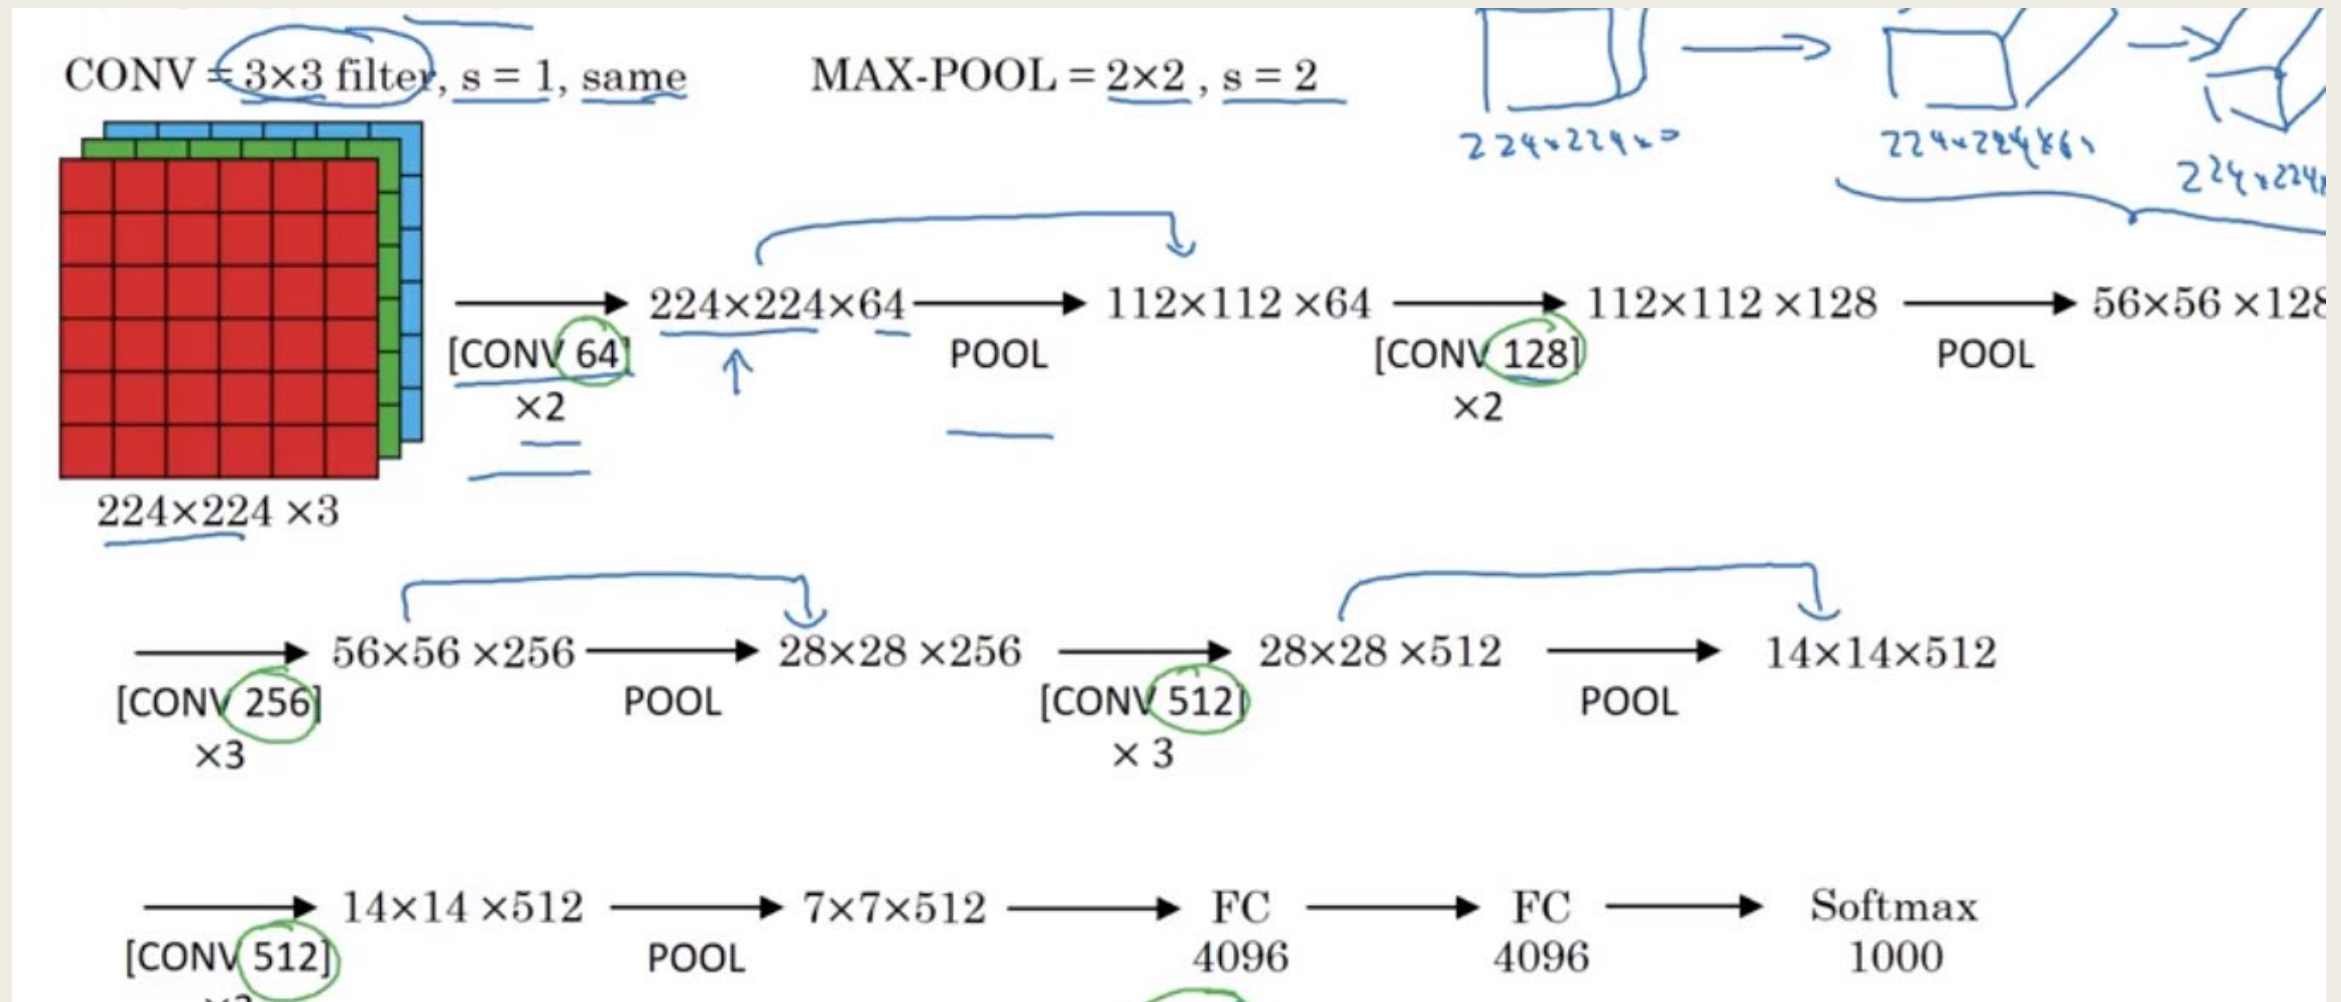


- Why convolutions?
 - **Parameter sharing:** A feature detector is probably useful in many parts of the image
 - **Sparsity of connections:** In each layer, each output value depends only on a small number of inputs


Examples: LeNet5 – AlexNet



Examples: VGG16



Residual Network






- Shortcut = Skip connection: $a^{[l+2]} = g(z^{[l+2]} + a^{[l]})$


- Residual network vs Plain network:

→ Identity function is easier for Residual block to learn (If $W^{[l+2]} = 0$ and $b^{[l+2]} = 0$)

Ps: If sizes don't match → $a^{[l+2]} = g(z^{[l+2]} + W_s, a^{[l]})$; Exp: $a^{[l+2]} \rightarrow 256, a^{[l]} \rightarrow 128, W_s \rightarrow (256, 128)$

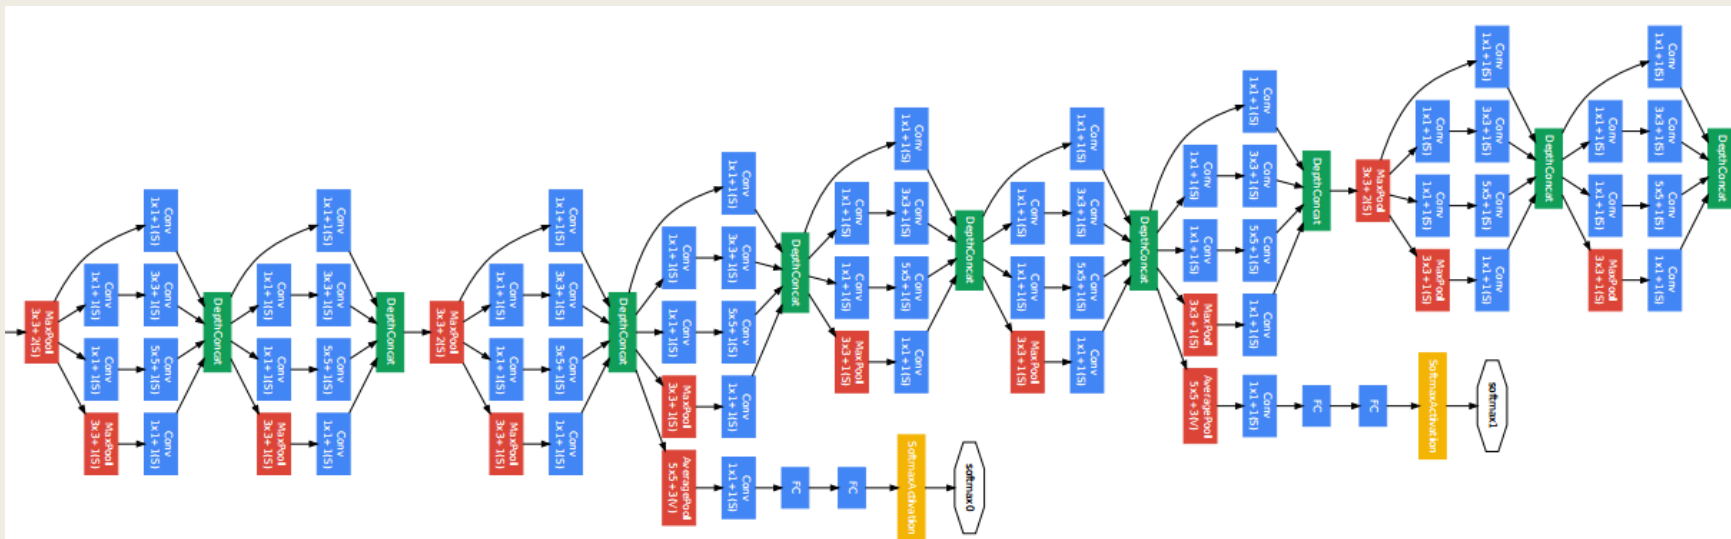
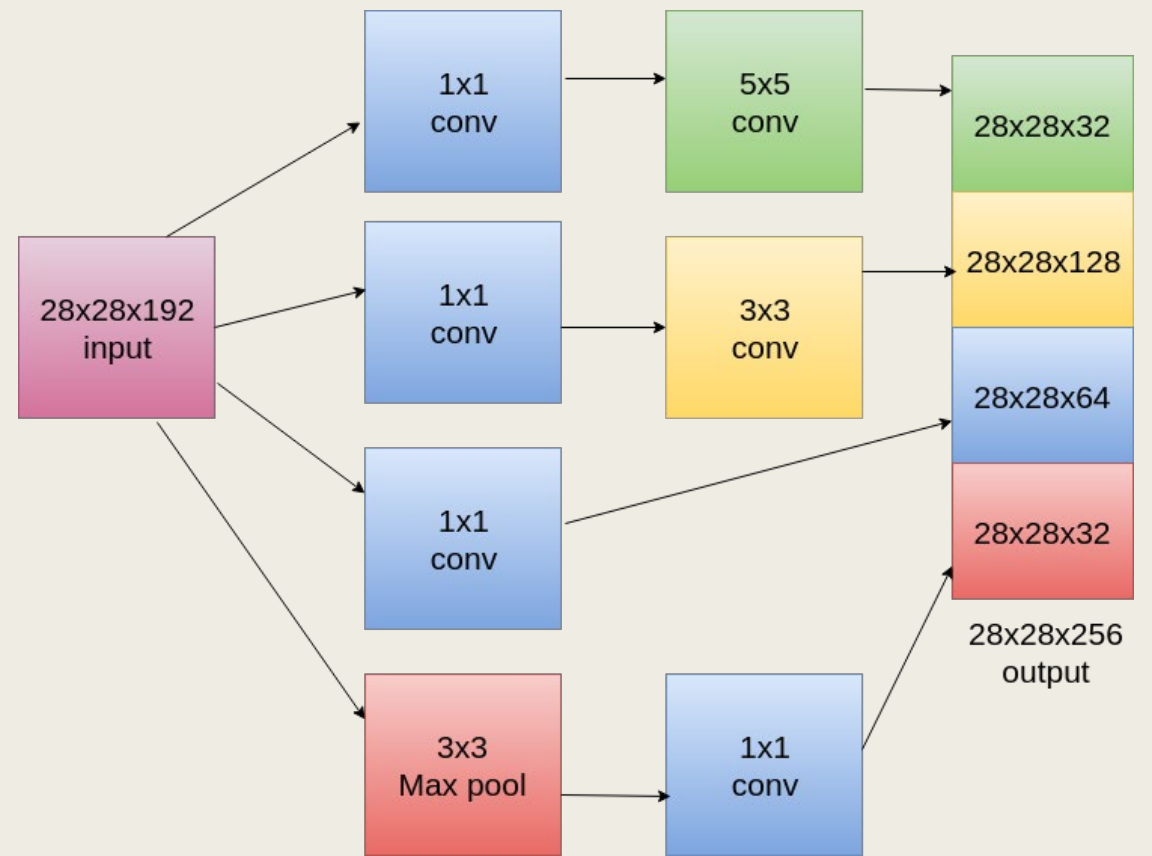


1x1 Convolution

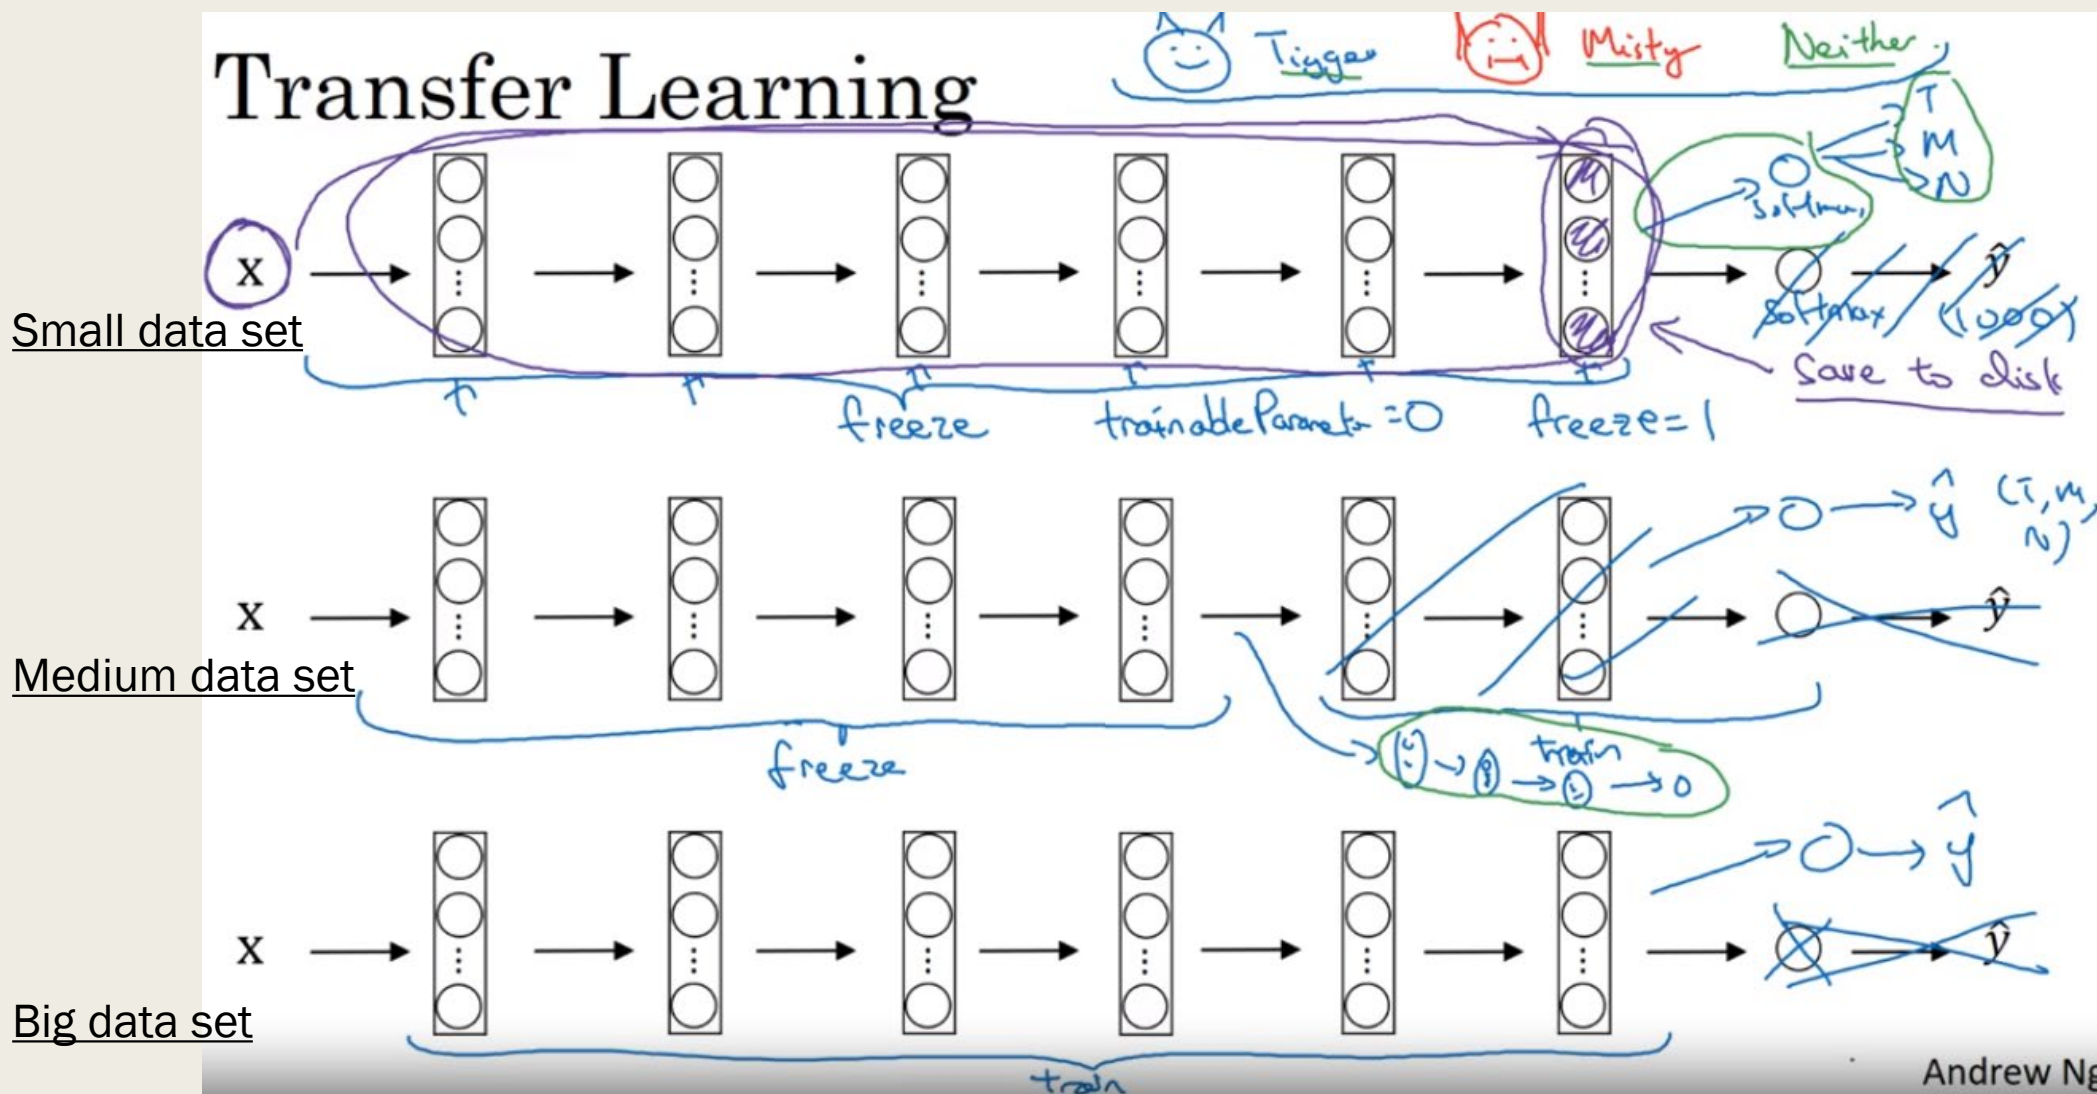


- **Goal:** Change size from $(n, n, value)$ to $(n, n, \#filters)$ thanks to $(1, 1, value)$ $\#filters$
- Reduce computational cost (1x1 layer is called « bottleneck » layer)
- **Tip:** Use open source code
 - Use architectures of networks published in the literature
 - Use open source implementations if possible
 - Use pretrained models and fine-tune on your dataset

Inception Network



Transfer Learning



Data Augmentation

- Techniques: Mirroring, Random cropping, Rotation, Shearing, Local warping, Color shifting...
- Implementing distortions during training → Parallel job using threads:
 - CPU thread to distort picture
 - CPU thread to train data
- Two sources of knowledge:
 - Labeled data
 - Hand engineered features/network architecture/other components
- **Little** data (More hand-engineering) vs **Much** data (Less hand-engineering ⇔ Simpler algorithms)
- **Little** data: Object detection → Image recognition → Speech recognition: **Much** data
- Tips for doing well on benchmarks:
 - **Ensembling**: Train several networks independently and average their outputs
 - **Multi-crop at test time**: Run classifier on multiple versions of test images and average results

10-crop



Keras {1}

- High level framework that provides additional abstractions (Higher than TensorFlow)
- In Keras, instead of creating a new variable on each step of forward propagation (X, Z1, A1, Z2, A2...) we just reassign X to a new value


```
def model(input_shape):  
    # Define the input placeholder as a tensor with shape input_shape. Think of this as your input image!  
    X_input = Input(input_shape)  
  
    # Zero-Padding: pads the border of X_input with zeroes  
    X = ZeroPadding2D((3, 3))(X_input)  
  
    # CONV -> BN -> RELU Block applied to X  
    X = Conv2D(32, (7, 7), strides = (1, 1), name = 'conv0')(X)  
    X = BatchNormalization(axis = 3, name = 'bn0')(X)  
    X = Activation('relu')(X)  
  
    # MAXPOOL  
    X = MaxPooling2D((2, 2), name='max_pool')(X)  
  
    # FLATTEN X (means convert it to a vector) + FULLYCONNECTED  
    X = Flatten()(X)  
    X = Dense(1, activation='sigmoid', name='fc')(X)  
  
    # Create model. This creates your Keras model instance, you'll use this instance to train/test the model.  
    model = Model(inputs = X_input, outputs = X, name='HappyModel')  
  
    return model
```

Keras {2}


1. Create the model by calling the function above
 2. Compile the model by calling `model.compile(optimizer="", loss="", metrics=["accuracy"])`
 3. Train the model on train data by calling `model.fit(X=,Y=,epochs=,batch_size=)`
 4. Test the model on test data by calling `model.evaluate(X=,Y=)`
- Useful tips:
- `model.summary()`: prints the details of your layers in a table with the sizes of its inputs/outputs
 - `plot.model()`: plots your graph in a nice layout. You can even save it as “.png” using SVG()

Classification, Localization, Detection


Image classification




Classification with
localization



Detection



Classification with Localization



- Output layer contains:
$$\begin{pmatrix} p_c \rightarrow 1 \text{ if object exist} \\ b_x \rightarrow \text{box } x_coordinate \\ b_y \rightarrow \text{box } y_coordinate \\ b_h \rightarrow \text{box height} \\ b_w \rightarrow \text{box width} \\ c_1 \rightarrow \text{1st object class} \\ c_2 \rightarrow \text{2nd object class} \\ c_3 \rightarrow \text{3rd object class} \end{pmatrix}; \text{ PS: If } p_c=0, \text{ we don't care about other values}$$


Landmark Detection & Non-Convolution Implementation of Sliding Windows

- Goal: Detect points (coordinates) instead of boxes (with height and width)


- Output layer contains:
$$\begin{pmatrix} p_c \rightarrow 1 \text{ if object exist} \\ l_{1x} \rightarrow \text{landmark1 } x_coordinate \\ l_{1y} \rightarrow \text{landmark1 } y_coordinate \\ l_{2x} \rightarrow \text{landmark2 } x_coordinate \\ l_{2y} \rightarrow \text{landmark2 } y_coordinate \\ \vdots \end{pmatrix}$$



Turning FC layer into Convolutional Layer




Convolution Implementation of Sliding Windows



- PS: Instead of working on each window by itself (non-convolution), we will work on all the windows together using convolutional implementation

YOLO Algorithm

- YOLO = You Only Look Once
- Output size $\rightarrow (\sqrt{\#grids}, \sqrt{\#grids}, 5 + \#classes)$, $5 = pc + b_x + b_y + b_h + b_w$
- PS: If $P_c=0$, other values don't matter
- PS1: $0 \leq b_x$ and $b_y \leq 1 \rightarrow$ Relative to the grid
- PS2: b_h and b_w can be > 1




Non-Max Suppression Algorithm

- Intersection over Union: $IoU = \frac{\text{size of intersection}}{\text{size of union}} \rightarrow \text{Correct if } IoU \geq 0,5$
 - IoU is measure of the overlap between two bounding boxes

- Algorithm:


1. Each output prediction is $\begin{pmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \end{pmatrix}$
2. Discard all boxes with $p_c \leq 0,6$
3. While there are any remaining boxes:
 1. Pick the box with the largest p_c and output that as a prediction
 2. Discard any remaining box with $IoU \geq 0,5$ with the box output in the previous step



Anchor Box Algorithm

- **Goal:** Find solution to overlapping objects


PS: With two anchor boxes, each object in training image is assigned to grid cell that contains object's midpoint and anchor box for the grid cell with highest IoU



- Output size $\rightarrow (\sqrt{\#grids}, \sqrt{\#grids}, (5 + \#classes) * \#AnchorBoxes), 5 = p_c + b_x + b_y + b_h + b_w$
- PS: If one $P_c=0$, other values with same anchor box don't matter




- **Drawbacks:**

- *It is bad when different objects have the same shape*
- *It is bad if 3 objects appear when having 2 anchor boxes*



Final Algorithm

1. For each grid cell, get 2 predicted bounding boxes
2. Get rid of low probability predictions
3. For each class (pedestrian, car, motorcycle) use non-max suppression to generate final predictions



Region Proposal: R-CNN

- It is a segmentation algorithm to reduce the calculations
 - *R-CNN: Propose regions. Classify proposed regions one at a time. Output level + bounding box*
 - *Fast R-CNN: Propose regions. Use convolution implementation of sliding windows to classify all the proposed regions*
 - *Faster R-CNN: Use convolutional network to propose regions*



One-shot Learning (1)

■ Face verification:

- Input image, name/ID
- Output whether the input image is that of the claimed person

■ Face recognition:

- Has a database of K persons
- Input image
- Output ID if the image is any of the K persons (or “not recognized”)

■ One-shot learning: Learning from one example to recognize the person again

→ Learning a “**similarity**” function: $d(img1, img2)$ = degree of difference between images

$$d(x^{(1)}, x^{(2)}) = \left\| f(x^{(1)}) - f(x^{(2)}) \right\|_2^2$$

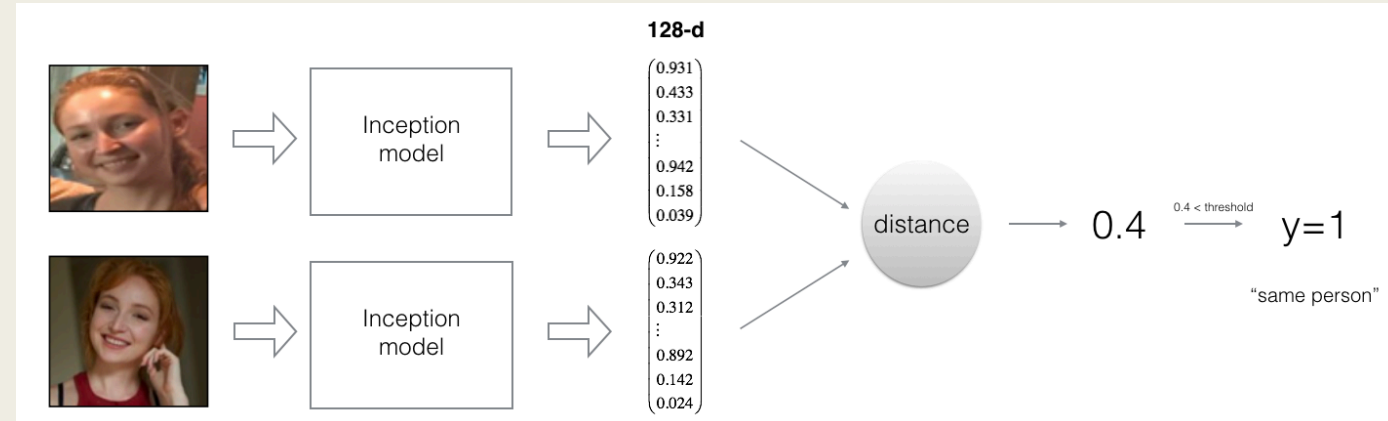
if $d(img1, img2) \leq \tau \rightarrow$ same

if $d(img1, img2) > \tau \rightarrow$ different

→ **Siamese Network**

- Person 1 → CNN → Softmax output (128 units) = encoding of person 1
- Person 2 → Same CNN → Softmax output = encoding of person 2
- Learn parameters so that if person 1/2 are the same person, d is small, and large otherwise

■ Famous Networks: FaceNet, DeepFace




One-shot Learning (2)

- We want $\|f(A) - f(P)\|^2 + \alpha \leq \|f(A) - f(N)\|^2 \Leftrightarrow d(A, P) + \alpha \leq d(A, N)$
 - *A: Anchor image, P: Positive, N: Negative, α : margin to make difference larger*
- Loss function: $L(A, P, N) = \max\left(\|f(A) - f(P)\|^2 - \|f(A) - f(N)\|^2 + \alpha, 0\right)$
- $J = \sum_{i=1}^n L(A^{(i)}, P^{(i)}, N^{(i)})$; Training set: 10k pictures of 1k persons for example
- **PS:** if A, P, N are chosen randomly, the inequality is easily satisfied
 - Choose triplets that are “hard” to train on $\Leftrightarrow d(A, P) \approx d(A, N)$
- **Learning the similarity function:** If a new person is added to the dataset, its output $f(x^{(\text{new})})$ will be compared to the already precomputed $f(x^{(i)})$ of the dataset

$$\hat{y} = \sigma \left(\sum_{k=1}^{128} w_i |f(x^{(i)})_k - f(x^{(j)})_k| + b \right)$$

Visualizing what a deep network is learning

- **Tip:** Pick a unit in layer I. Find the nine image patches that maximize the unit's activation.



Layer 1



Layer 2



Layer 3




Layer 4



Layer 5

Neural Style Transfer



- $J(G) = \alpha J_{content}(C, G) + \beta J_{style}(S, G)$

- Algorithm:

1. Initialize G randomly

2. Use gradient descent to minimize $J(G)$: $G := G - \frac{d}{dG} J(G)$

- $J_{content}^{[l]}(C, G) = \frac{1}{2} \left\| a^{[l](C)} - a^{[l](G)} \right\|^2$

→ The difference between the activation of the two layers. We want them to be similar

- **Style** = Correlation between activations across channels

- **Style matrix**: = **Gram matrix**: $G_{kk'}^{[l]} = \sum_{i=1}^{n_H^{[l]}} \sum_{j=1}^{n_W^{[l]}} a_{ijk}^{[l]} a_{ijk'}^{[l]}$

- $J_{style}^{[l]}(S, G) = \frac{1}{(2n_H^{[l]}n_W^{[l]}n_C^{[l]})^2} \sum_{k=1}^{n_C^{[l]}} \sum_{k'=1}^{n_C^{[l]}} (G_{kk'}^{[l](S)} - G_{kk'}^{[l](G)})^2 = \frac{1}{(2n_H^{[l]}n_W^{[l]}n_C^{[l]})^2} \left\| G^{[l](S)} - G^{[l](G)} \right\|_F^2$

- $J_{style}(S, G) = \sum_l \lambda^{[l]} \cdot J_{style}^{[l]}(S, G)$

